

Software security lab01

Davide Caria
Tee Wei Teoh
Sara Lehne Engesvik

11 October 2023

1 Code analysis of exercise n2: WinLoose

1.1 Point 1

The program has been compiled with the following command:

```
gcc -fno-stack-protector -o winloose winloose.c
```

and tested with different inputs for the *./winloose* file. As an example we tried with several inputs and wrote down the outputs.

- *./winloose* 3 3
- *./winloose* 10 10
- *./winloose* 3 18
- *./winloose* 30 7
- *./winloose* 5 8
- *./winloose* 0 11
- *./winloose* 1 8

1.2 Point 2

According to the previous list of possible inputs we deduced the following:

- $argv[2] > 8$ the program crashes into a segmentation fault error
- $argv[2] < 8$ the output is always *Youloose...*
- $argv[2] = 8$ and $argv[1] = 0$ the output is *Youloose...*
- $argv[2] = 8$ and $argv[1] \neq 0$ the output is always *YouWin!*

The guessing of the numbers suggest that something makes the program crash when the second argument is above 8, looking at the C code we can see that this is due to a buffer overflow in line 15. If the program does not crash or loop endlessly it can produce two outputs which depend on the value of the variable *X*. Specifically, there are no boundary checks of the input given by the user. Therefore, if we use a value that exceeds the array length we step into the *X* variable location and overwrite it. This means that if the first argument is different from zero, in the last iteration it will become the new value of *X* making us win the game.

1.3 Point 3

To confirm our guesses we disassembled the code using the *objdump* tool. The specific disassembled code on the lab machine was the following:

```
0000000000001169 <main>:
1169: f3 0f 1e fa      endbr64
116d: 55              push %rbp
116e: 48 89 e5         mov %rsp,%rbp
1171: 48 83 ec 20      sub $0x20,%rsp
1175: 89 7d ec         mov %edi,-0x14(%rbp)
1178: 48 89 75 e0      mov %rsi,-0x20(%rbp)
117c: c6 45 fb 00     movb $0x0,-0x5(%rbp)
1180: c7 45 fc 00 00 00 00 movl $0x0,-0x4(%rbp)
1187: eb 22           jmp 11ab <main+0x42>
1189: 48 8b 45 e0      mov -0x20(%rbp),%rax
118d: 48 83 c0 08      add $0x8,%rax
1191: 48 8b 00         mov (%rax),%rax
1194: 48 89 c7         mov %rax,%rdi
1197: e8 d4 fe ff ff  call 1070 <atoi@plt>
119c: 89 c2           mov %eax,%edx
119e: 8b 45 fc         mov -0x4(%rbp),%eax
11a1: 48 98           cltq
11a3: 88 54 05 f3     mov %dl,-0xd(%rbp,%rax,1)
11a7: 83 45 fc 01     addl $0x1,-0x4(%rbp)
11ab: 48 8b 45 e0      mov -0x20(%rbp),%rax
11af: 48 83 c0 10      add $0x10,%rax
11b3: 48 8b 00         mov (%rax),%rax
11b6: 48 89 c7         mov %rax,%rdi
11b9: e8 b2 fe ff ff  call 1070 <atoi@plt>
11be: 39 45 fc         cmp %eax,-0x4(%rbp)
11c1: 7e c6           jle 11b9 <main+0x20>
11c3: 80 7d fb 00     cmpl $0x0,-0x5(%rbp)
11c7: 74 11           je 11da <main+0x71>
11c9: 48 8d 05 34 0e 00 00 lea 0xe34(%rip),%rax # 2004 <_IO_stdin_used+0x4>
11d0: 48 89 c7         mov %rax,%rdi
11d3: e8 88 fe ff ff  call 1060 <puts@plt>
11d8: eb 0f           jmp 11e9 <main+0x08>
11da: 48 8d 05 2d 0e 00 00 lea 0xe2d(%rip),%rax # 200e <_IO_stdin_used+0xe>
11e1: 48 89 c7         mov %rax,%rdi
11e4: e8 77 fe ff ff  call 1060 <puts@plt>
11e9: b8 00 00 00 00  mov $0x0,%eax
11ee: c9             leave
11ef: c3             ret
```

Figure 1: Assembly code for *winloose* using *objdump*

Then, after analyzing the output of *objdump* we worked out a possible memory stack structure, including the location of the variables and arrays. We assumed the following configuration:

1.5 Point 6

After disassembling the protected code with *objdump* tool we obtained the following assembly instructions:

```
0000000000001189 <main>:
1189: f3 0f 1e fa      endbr64
118d: 55              push    rbp
118e: 48 89 e5        mov     rbp, rsp
1191: 48 83 ec 30      sub     rsp, 0x30
1195: 89 7d dc        mov     DWORD PTR [rbp-0x24], edi
1198: 48 89 75 d0      mov     QWORD PTR [rbp-0x30], rsi
119c: 64 48 8b 04 25 28 00 mov     rax, QWORD PTR fs:0x28
11a3: 00 00          mov     QWORD PTR [rbp-0x8], rax
11a5: 48 89 45 f8      mov     eax, eax
11a9: 31 c0          xor     eax, eax
11ab: c6 45 eb 00      mov     BYTE PTR [rbp-0x15], 0x0
11af: c7 45 ec 00 00 00 00 mov     DWORD PTR [rbp-0x14], 0x0
11b6: eb 22          jnp     11da <main+0x51>
11b8: 48 8b 45 d0      mov     rax, QWORD PTR [rbp-0x30]
11bc: 48 83 c0 08      add     rax, 0x8
11c0: 48 8b 00        mov     rax, QWORD PTR [rax]
11c3: 48 89 c7        mov     rdi, rax
11c6: e8 c5 fe ff ff  call    1090 <atoi@plt>
11cb: 89 c2          mov     edx, eax
11cd: 8b 45 ec        mov     eax, DWORD PTR [rbp-0x14]
11d0: 48 98          cdq     eax
11d2: 88 54 05 f0      mov     BYTE PTR [rbp+rax*1-0x10], dl
11d6: 83 45 ec 01      add     DWORD PTR [rbp-0x14], 0x1
11da: 48 8b 45 d0      mov     rax, QWORD PTR [rbp-0x30]
11de: 48 83 c0 10      add     rax, 0x10
11e2: 48 8b 00        mov     rax, QWORD PTR [rax]
11e5: 48 89 c7        mov     rdi, rax
11e8: e8 a3 fe ff ff  call    1090 <atoi@plt>
11ed: 39 45 ec        cmp     DWORD PTR [rbp-0x14], eax
11f0: 7e c6          jle     11b8 <main+0x2f>
11f2: 89 7d eb 00      cmp     BYTE PTR [rbp-0x15], 0x0
11f6: 74 11          je      1209 <main+0x88>
11f8: 48 8d 05 05 0e 00 00 lea     rax, [rip+0xe05] # 2004 <_IO_stdin_used+0x4>
11ff: 48 89 c7        mov     rdi, rax
1202: e8 69 fe ff ff  call    1070 <puts@plt>
1207: eb 0f          jmp     1218 <main+0x8f>
1209: 48 8d 85 fe 0d 00 00 lea     rax, [rip+0xdfe] # 200e <_IO_stdin_used+0xe>
1210: 48 89 c7        mov     rdi, rax
1213: e8 58 fe ff ff  call    1070 <puts@plt>
1218: b8 00 00 00 00  mov     eax, 0x0
121d: 48 8b 55 f8      mov     rdx, QWORD PTR [rbp-0x8]
1221: 64 48 2b 14 25 28 00 sub     rdx, QWORD PTR fs:0x28
1228: 00 00          je      1231 <main+0xa8>
122a: 74 05          je      1231 <main+0xa8>
122c: e8 4f fe ff ff  call    1080 <__stack_chk_fail@plt>
1231: c9            leave   %rax
1232: c3            ret
```

Figure 3: Assembly code for *winloose* with stack protection

We can see that there are some differences in how the stack is protected in this new version of the assembly code: First there has been given more space to the stack and the buffers that are allocated. Then the compiler inserted more *cmp* statements to check if we are exceeding the boundaries. Specifically, a canary value is initialized by the compiler to a random value and placed where the buffer overflow can happen. At the end of the loop the canary value is checked again and if it is different with respect to the original value the program jumps to a routine. The routine is *stack_chk_fail* and it handles the buffer overflow that just occurred.

2 Code analysis of exercise n4: ExecShell

2.1 Point 1

After analyzing the the provided c code, we found three vulnerabilities:

- No input sanitization
- Buffer-overflow

- Use-after-free

2.1.1 No input sanitization

For what concerns input-sanitization, we can see that there is no protection or check that prevents the insertion of wrong values. Assuming that a correct user is trying to operate the tool, he has no boundaries neither on the size nor in the content. This means, as example, that a string that does not represent a directory is considered valid.

2.1.2 Buffer-overflow

There is an attempt to limit the possibility of buffer-overflow at line 22:

```
if (strlen(a) > 14) { ... }
```

Nevertheless, this does not solve the problem completely and leaves the chance of buffer-overflow later in the code. Specifically, at line 36 we can see that a new buffer is allocated with slightly more space compared to the first allocation.

```
p3 = (char *) malloc (24 * sizeof (char)); // allocation 2
```

This may help with some standard use cases, for example when there are small paths as inputs. However, any input longer than 24 characters will result in an overflow.

2.1.3 User-after-free

This last vulnerability presents itself when we have an input larger than 16 character. In line 24 the program frees the pointer to the buffer *p*:

```
free (p);
```

and later in the code the pointer *p* is used again as argument for the system call *system()*. This is the perfect example of use-after-free vulnerability. In this particular case, the vulnerability is combined with wrong code logic: the pointer *p* refers only to the original buffer so even if the user insert a valid string during the second trial, the program will still execute:

```
system(p)
```

ignoring the new allocated string *p3*.

2.2 Point 2

A possible way to execute an arbitrary shell command can leverage the character ; which concatenates multiple command. As an example we could provide a valid path for the *ls* tool but concatenate a malicious code.

```
input: /tmp ; rm -r /
```

This input will be concatenated with the predefined string in the program to form a valid shell command:

```
$ ls /tmp ; rm -r /
```

This command will list all the elements under */tmp* directory and subsequently it will try to wipe the disk from the root folder */*. Most probably the command will not work thanks to the *admin* write restriction. Nevertheless, other sub-folder could be deleted if the read-write policy is not properly configured.

3 Sub-rule explanation

The team choose *Rule 08* and sub-rule *MEM30 – C* from the CERT Secure Coding Standard.

RULE	SEVERITY	LIKELIHOOD	REMEDIATION COST
MEM30-C	High	Likely	Medium

The given rule, MEM30-C, focuses on avoiding access to freed memory, which can lead to undefined behavior and exploitable vulnerabilities. It outlines the risks associated with accessing or modifying memory that has been deallocated using functions like *free()* or *realloc()*. The main problems addressed by the rule include reading or writing to memory after it has been freed and the potential for double-free vulnerabilities

3.1 Problem 1: Accessing Freed Memory

Accessing a pointer to memory that has been deallocated (dangling pointer) is undefined behavior and can lead to exploitable vulnerabilities. Reading or writing to this memory can result in abnormal program termination and denial-of-service attacks

3.2 Problem 2: Double-Free Vulnerabilities

Freeing memory multiple times, either directly or through incorrect usage of *realloc()*, can cause double-free vulnerabilities. Double-free can lead to the execution of arbitrary code with the permissions of the vulnerable process.

An example of **non compliant code**:

```
#include <stdlib.h>

void f(char *c_str1, size_t size) {
    char *c_str2 = (char *)realloc(c_str1, size);
    if (c_str2 == NULL) {
        free(c_str1);
    }
}
```

An example of **compliant code**:

```
#include <stdlib.h>

void f(char *c_str1, size_t size) {
    if (size != 0) {
        char *c_str2 = (char *)realloc(c_str1, size);
        if (c_str2 == NULL) {
            free(c_str1);
        }
    }
    else {
        free(c_str1);
    }
}
```

In the non-compliant code, the *realloc()* function is called with a size argument of zero, which can lead to potential issues. According to the C standard, when *realloc()* is called with a size of zero, it may behave like *free()* or return a unique pointer that can be used for future allocations. If the returned pointer is used and the original pointer (*c_str1*) is already de-allocated, this can lead to undefined behavior. In the compliant code, a check for *size != 0* is added before calling *realloc()*. This ensures that *realloc()* is only called when the size is non-zero, addressing the issue and preventing potential undefined behavior associated with calling *realloc()* with a size of zero.

3.3 Solution for problem 1 (Accessing Freed Memory)

Ensure that memory is not accessed (read or written) after it has been freed. Properly manage the order of operations to prevent such access. Free the memory after its final use to avoid any potential access to the freed memory

3.4 Solution for problem 2 (Double-Free Vulnerabilities)

Properly handle memory deallocation to prevent double-free vulnerabilities. Ensure that memory is freed only once and is not attempted to be freed again